

Machine Learning Models for Walmart Inventory Forecasting

This document details the machine learning models implemented for forecasting Units Sold in the Walmart inventory management system. We've utilized three distinct models—XGBoost, LSTM, and Prophet—each chosen for its strengths in handling different aspects of time-series data, and then combined them into an ensemble.

1. Data Preparation and Feature Engineering

Before model training, the raw `retail_inventory_forecast.csv` dataset underwent a crucial data preparation and feature engineering phase. This involved:

- **Cleaning Missing Values:**
 - Missing Discount values were filled with 0.
 - Missing Competitor Pricing values were filled with the mean of the existing values.
 - Other numerical columns (Inventory Level, Units Ordered, Demand Forecast, Price) were filled with 0 if missing.
- **Feature Creation from Date:**
 - Date column was converted to datetime objects.
 - New time-based features were extracted: `day_of_week`, `week_of_year`, month, and year.
- **Lag Features for Units Sold:**
 - To capture temporal dependencies, lag features for Units Sold were created by shifting values: `units_sold_lag_1` (previous day's sales), `units_sold_lag_7` (previous week's sales), and `units_sold_lag_30` (previous month's sales). These were calculated per unique Store ID and Product ID.
- **Rolling Averages for Units Sold:**
 - To smooth out short-term fluctuations and capture trends, rolling averages for Units Sold were computed: `units_sold_rolling_avg_7` (7-day moving average) and `units_sold_rolling_avg_30` (30-day moving average). These were also calculated per unique Store ID and Product ID.
- **Column Standardization:** Critical column names like Holiday/Promotion and Competitor Pricing were explicitly standardized to ensure consistency across all scripts.

The processed data was then saved to `retail_inventory_forecast_processed.csv` and loaded into a dedicated MongoDB Atlas collection (`forecast_data_processed`).

2. Core Models and Their Outputs

We trained three distinct models, each with its own strengths. The models were

evaluated on a validation set (the last 20% of the chronologically sorted data) using the Mean Absolute Error (MAE) metric.

2.1. XGBoost Regressor

- **Best For:** Products whose sales are heavily influenced by promotions, price changes, competitor activity, and other structured external factors. It excels at capturing complex non-linear relationships between features and the target.
- **Implementation:**
 - **Input Features:** Price, Discount, Holiday/Promotion, Competitor Pricing, day_of_week, month, year, units_sold_lag_1, units_sold_lag_7, units_sold_lag_30, units_sold_rolling_avg_7, units_sold_rolling_avg_30.
 - **Target Variable:** Units Sold.
 - **Training:** The model was trained using a time-series split (first 80% for training, last 20% for validation). Due to the specific XGBoost version, direct early_stopping_rounds was used in the fit method.
- **Output (Validation MAE):**
XGBoost Validation MAE: 80.06

This indicates that, on average, the XGBoost model's predictions for Units Sold on the validation set were off by approximately 80.06 units.

2.2. LSTM (Long Short-Term Memory) Network

- **Best For:** Capturing complex temporal patterns and long-term dependencies in sales data. LSTMs are a type of recurrent neural network particularly suited for sequence prediction tasks.
- **Implementation:**
 - **Input Features:** Sequences of past Units Sold, day_of_week, and month. A LOOK_BACK period of 7 (i.e., using the past 7 days' data) was used to create sequences.
 - **Data Preparation:** Features and the target (Units Sold) were scaled using MinMaxScaler to a range of (0, 1). Data was grouped by Store ID and Product ID to create sequences for each unique time series.
 - **Architecture:** A sequential Keras model with two LSTM layers (50 units each), followed by Dropout layers (0.2), and a final Dense layer (1 unit) for the prediction.
 - **Training:** Trained with adam optimizer and mean_squared_error loss. Early stopping and model checkpointing were used to save the best model based on validation loss.
- **Output (Validation MAE):**

LSTM Validation Mean Absolute Error (MAE): 88.97

The LSTM model's predictions were, on average, off by about 88.97 units on the validation set.

2.3. Prophet Model

- **Best For:** Products with strong, predictable seasonality (e.g., weekly, yearly) and recurring holiday effects. It's designed for business forecasting and handles missing data, outliers, and trend changes automatically.
- **Implementation:**
 - **Input Features:** Prophet primarily requires a ds (date) column and a y (target) column. For this implementation, Units Sold was aggregated to a *daily total* (y) and Date became ds.
 - **Holiday/Regressor Handling:** Holiday/Promotion was incorporated as a custom holiday effect. Competitor Pricing was added as an extra_regressor, with its daily mean used for future dates.
 - **Training:** The model automatically decomposes the time series into trend, seasonality (yearly, weekly, daily), and holiday components.
- **Output (Validation MAE):**

Prophet Validation Mean Absolute Error (MAE): 850.00

The Prophet model, when predicting daily total Units Sold, had an average error of approximately 850.00 units on the validation set. This higher MAE compared to XGBoost and LSTM is expected because Prophet is forecasting the *sum* of units sold across all stores/products, which is a much larger number, whereas XGBoost and LSTM were implicitly predicting per-item-per-day.

3. Ensemble Creation and Weighted Averaging

The ensemble combines the strengths of the individual models to produce a more robust and accurate final forecast.

- **Process:**
 1. **Load Models:** The trained XGBoost, LSTM, and Prophet models were loaded.
 2. **Generate Predictions:** Each model generated predictions on the shared validation dataset.
 3. **Calculate Individual MAEs:** The Mean Absolute Error was calculated for each model's predictions against the actual values on the validation set.
 4. **Determine Weights:** Weights for each model were assigned inversely proportional to their validation MAE. This means a model with a lower MAE (better performance) receives a higher weight in the final ensemble.

- $w_i = \text{frac1/MAE_i} / \sum(1/\text{MAE}_j)$
 - If an MAE was 0 (perfect prediction, which is rare), its inverse was set to 0 to avoid division by zero, or a fallback to equal weights was used if all MAEs were problematic.
5. **Combine Predictions (Weighted Average):** The final ensemble forecast was calculated as a weighted average of the individual model predictions.
- **Granularity Note:** Due to the differing prediction granularities (XGBoost/LSTM: per-item-per-day; Prophet: daily total), a simplification was made. XGBoost's per-item predictions were aggregated to daily totals to align with Prophet. **LSTM's predictions were not directly summed into the final ensemble forecast due to the complexity of aligning its sequence-based outputs to a daily total in a straightforward manner.** However, LSTM's calculated weight still influenced the overall balance of the ensemble by reducing the relative weights of XGBoost and Prophet. The ensemble specifically combined the daily aggregated XGBoost and Prophet predictions.
- **Output (Ensemble Weights and MAE):**
 Calculated Ensemble Weights:
 XGBoost Weight: 0.5015
 LSTM Weight: 0.4513
 Prophet Weight: 0.0472
 Sum of Weights: 1.0000

Ensemble (XGBoost + Prophet) Validation MAE: 801.04

The ensemble weights indicate that XGBoost and LSTM contributed the most to the overall weighting scheme, reflecting their lower individual MAEs. Prophet, despite its higher MAE for daily totals, still contributed a small weight.

The final ensemble (combining XGBoost and Prophet for daily totals) achieved an MAE of **801.04**. This is slightly better than Prophet's individual MAE (850.00) for daily totals, demonstrating the benefit of combining models.

Summary

We have successfully built and evaluated three machine learning models for demand forecasting and created a weighted ensemble. The data preprocessing step was crucial for preparing the data for each model's specific requirements. The ensemble approach allows us to leverage the strengths of different modeling techniques to achieve a more robust and potentially more accurate overall forecast for Walmart's

inventory.

The next steps would involve integrating these trained models into the Flask backend to provide real-time forecasting capabilities to the React frontend.